Assignment 1

Filipe Pires [85122], João Alegria [85048]

Information Retrieval

Department of Electronics, Telecommunications and Informatics

University of Aveiro

October 2, 2019

Introduction

This report aims to describe the work developed for the first assignment of the discipline of 'Information Retrieval', explaining the overall processing pipeline. We include a short description of each class developed and of the respective methods, as well as the instructions on how to run our code.

The program implemented in Python version 3 has the purpose of indexing documents given as input in a compressed format. This document indexer consists of a corpus reader, a tokenizer and the actual indexer.

Along with the description of the solution, we also answer to a few questions proposed for the assignment (*1*).

1. Architecture

In order to maintain a modular architecture, we resorted to the Python library *Abstract Base Classes* (or simply ABC) (*2*). This module provided us with the infrastructure for defining abstract classes for each of our own modules. As seen in Fig.1, the 3 base classes all derive from an ABC.

The choice of this form of architecture was due to several reasons: modules allow the reduced coupling between system components, making it easy to replace or add new ones; creating variations of the solution (as we will see on the `Tokenizer` class) becomes far simpler and easier to manage; adopting this program structure also helps making the extension/adaptability of the program to other corpora structures easier to implement.
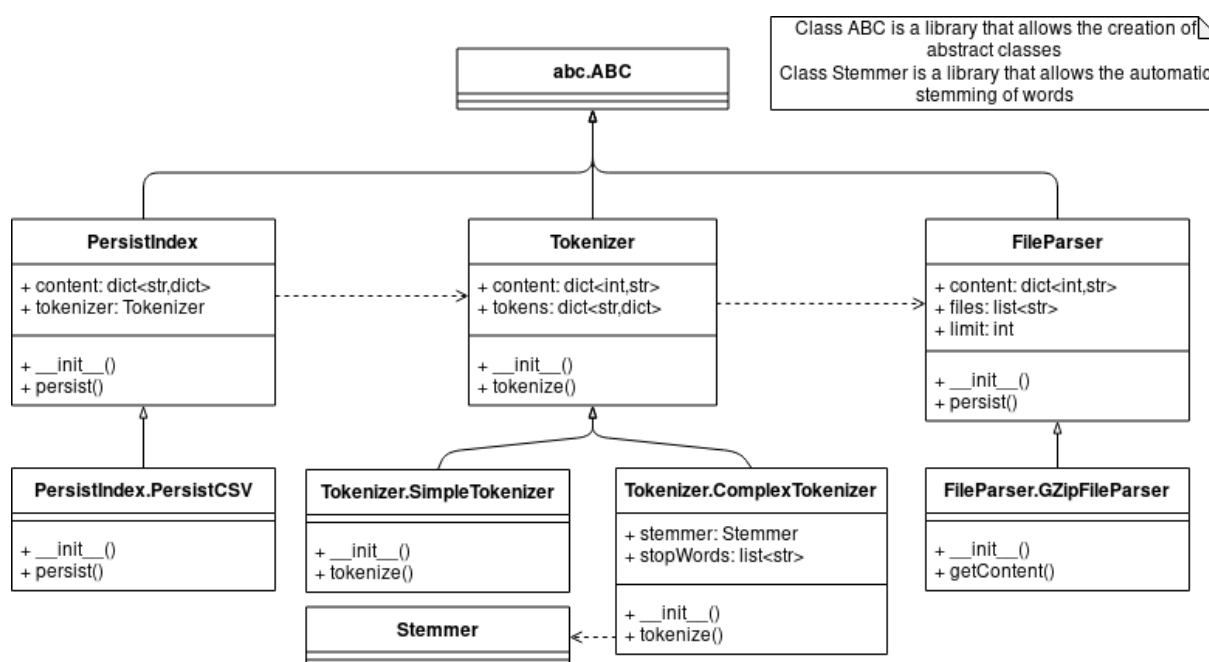

Figure 1: Program's class diagram.

The class `FileParser` is an abstract class for all file parsers. In our case, we only needed to develop 1 derived class to parse files on .gz format. It is `GZipFileParser` that actually reads the various compressed files passed to the program and dinamically processes them, meaning that the program decompresses the file line by line to retrieve the text bytes and then creates a dictionary that makes direct relation between a document id("PMID") and it's title("TI"). To improve our reading performance we used a reading buffer that significantly decreases the execution time.

`Tokenizer` is the abstract base class for our 2 Tokenizer implementations. This class depends on the `FileParser` as it uses the dictionary created by the File Parser. The `SimpleTokenizer` contains only one method called *tokenize()*, defined by the parent, that replaces all non-alphabetic characters by a space, lowercases the entire content, splits on whitespace, and ignores all tokens (all resulting strings after the split operation) with less than 3 characters. The `ComplexTokenizer`, on the other hand, is, as the name indicates, a more complex implementation aimed to return a more valuable set of tokens. This class integrates a stopword filter to ignore words considered not relevant and a stemmer to standardize all words, i.e. it normalizes the terms (or tokens) according to an algorithm that considers the English language by conflating words with the same linguistic base form. The stemmer used was PyStemmer (*3*), which is an implementation of Porter Stemmer, as recommended on the assignment instructions. For this reason, our Complex Tokenizer depends on the class `Stemmer`. In addition this tokenizer also implements custom rules conserning puctuation and more. We opted to implement restrictions in a way that maintains dates, emails, phone numbers, money amounts and numbers in general, but eliminates any kind of regular punctuation and separates words connected with hyphens or underscores, considering the resulting terms as tokens as well.

Finally we have the `PersistIndex` abstract class. Our implementation, `PersistCSV`, is the class responsible for persisting the ids of the documents where a given token appears and the number of occurrences in each of those documents, for all of the tokens from the entire vocabulary of the input files. This is done by storing the data present in a dictionary constructed by one of the previous tokenizers in a text file using the method *persist()*. To try to improve our performance we adopted an batch processing like system, were we store the token, all the document ID's where it appeared and the numbers of occurences on each document at the same time. After testing we observed that this approach improve our performance and execution time.

2. Data Flow

The file `Indexer.py` serves as the command to be executed in order to indexing all documents passed as arguments. This command has the following format:

```
$ python3 Indexer.py [-h] [-o outputFile] [-l limit] \\
    [-t tokenizer] inputFile1 [inputFile2]+
```

Here, `-h` is the option that presents the manual for the command usage. Options `-o` and `-l` allow the definition of the output file's name and of the limit for the number of lines to be processed in each input file. Option `-t` makes possible for the user to choose the type of tokenizer to be used, and the alternatives are: 'simple' for the use of the `SimpleTokenizer` class, and 'complex' for the `ComplexTokenizer` class. The previous arguments are all

optional and the actual values for these arguments must appear right after the respective options. The final argument(s) of the command must be the name(s) of the input file(s) to be indexed.

Once the arguments are validated, the program passes the file name(s) to the File Parser, already capable of processing multiple files. Once the parsing process is completed, the chosen Tokenizer accesses the data and tokenizes everything, that in this case is the title of each document. Lastly, the information created by the tokenizer - the actual index - is persisted into an output file, making possible its consultation and analysis.

Below we present both implementations of the method *tokenize()* for comparison purposes.

```python
def tokenize(self): # of class SimpleTokenizer
  for docID, docContent in self.content.items():
      token = re.sub("[^a-z]+", " ", docContent.lower()).split()
      for t in token:
          if len(t) >= 3:
              if t not in self.tokens:
                  self.tokens[t] = {docID: 1}
              elif docID not in self.tokens[t]:
                  self.tokens[t][docID] = 1
              else:
                  self.tokens[t][docID] = self.tokens[t][docID]+1
  return self.tokens

def tokenize(self): # of class ComplexTokenizer
  super().tokenize()
  for docID, docContent in self.content.items():
      normalizeData = docContent.lower()
      token = re.split(" +", normalizeData)
      stemmedTokens = self.stemmer.stemWords(token)
      for t in stemmedTokens:
          t = t if re.match(
              "[0-9]*(/|-)[0-9]*(/|-)[0-9]*", t) else re.sub(
              "([,;.:?!\(\)\[\]\{\}/\"\|])", " ", t)
          additionalWords = re.split(" +|_+|-+", t)
          t = additionalWords[0]
          tokens = stemmedTokens + \
              self.stemmer.stemWords(additionalWords[1:])
          if t not in self.stopWords:
              if t not in self.tokens:
                  self.tokens[t] = {docID: 1}
              elif docID not in self.tokens[t]:
                  self.tokens[t][docID] = 1
              else:
                  self.tokens[t][docID] = self.tokens[t][docID]+1
  return self.tokens
```

3. Discussion

To test the capabilities of the developed software, we indexed 2 large compressed files made available along with the assignment description with the names `2004_TREC_ASCII_MEDLINE_1.gz` and `2004_TREC_ASCII_MEDLINE_2.gz`. Each file is a collection (corpus) of documents. In this chapter we discuss our implementation's performance over these files and answer the following questions:

a) What is the total indexing time and final index size on disk?

b) What is the vocabulary size?

c) What are the ten first terms (in alphabetic order) that appear in only one document (document frequency = 1)?

d) What are the ten terms with highest document frequency?

To answer question a), we simple added the command `time` in the beginning of the execution of the `Indexer.py` (e.g. `$time python3 Indexer.py (...)`). This command returned us 3 measured times during the execution, but we were interested only in one of them - the Elapsed real time (in seconds) since the execution start. After all code optimizations, the best value we could achieve was of `5'48''` (five minutes and forty-eight seconds) and an index size on disk of `470.3 Mbs` with the `SimpleTokenizer` and of `10'52''` (ten minutes and fifty-two seconds) and an index size on disk of `406.1 Mbs` with the `ComplexTokenizer`.

For the remaining questions, we developed a small script called `IndexAnalyzer.py` that processes the output file passed as argument. The calculated size of the vocabulary generated by the Simple Tokenizer was `346221` unique terms and of the vocabulary generated by the Complex Tokenizer was `N`.

The 10 first terms (ordered alphabetically) with document frequency of 1 of the Simple and Complex Tokenizers are the following:

```
Simple:
'aaaat', 'aaab', 'aaac', 'aaacr', 'aaact', 'aaaction', 'aaad',
'aaaga', 'aaagat', 'aaah'

Complex:
.................
```

These were retrieved first filtering out all terms with document frequency above 1 and then ordering the remaining tokens alphabetically with the use of the *quicksort* algorithm.

The 10 terms with highest document frequency (for both tokenizers) are the following:

```
Simple:
Terms: 'and', 'the', 'with', 'for', 'from', 'patients', 'human',
'cell', 'cells', 'study'
Total Number of Occurences: 2044099

Complex:
Terms:
Total Number of Occurences:
```

To retrieve these terms, we stored the maximum number of occurrences during the output file's reading process and then search for the terms with that document frequency; if these terms are less than 10, we proceed on filtering for documents with the following highest frequency and repeat the process until the 10 terms are found.

Let us now take a closer look at these results. Starting with the indexing time, the indicated values were the minimum achievable after optimizing our code as much as we could. These execution times were reduced mainly due to the use of a Buffer Reader and of batch writings to the output file. The achieved sizes on disk of the indexes were the outcome of our way of treating the data. By balancing the data normalization and reducing the lost of knowledge throughout the processing pipeline, our solution seems to deliver indexes of considerably good value. This space is smaller for the Complex Tokenizer as it is able to remove unnecessary characters and compress the data in a smarter way.

The reason for the answers to questions b) and c) to be different depending on the tokenizer used is simple to understand by looking at their implementation. For example, as the ComplexTokenizer splits tokens by hyphens, this alone will already increase the vocabulary for the respective index. Another example with this tokenizer is that we decided it should not eliminate some symbols from words and, when these symbols occur in the beggining of a word, they are sorted as words that appear before the letter 'a', altering, in consequence, the answer to question c).

Again, for the answers to the final question, the tokenizer implementations alter the results as they treat specific words in different ways. For example, the ComplexTokenizer eliminates stopwords such as 'and' that occur massively in most English texts.

Conclusions

After completing the assignment, we drew a few conclusions regarding our solutions and the whole concept of indexing files containing textual information.

First of all, we had to learn how to deal with the actual tokenization of words and what rules should be taken in consideration when normalizing them. In this sense, the delivered code is the result of this learning process.

Second, a few unpredicted setbacks such as the extensive time that Python took when reading the input files without a buffer in the middle, made us realize more clearly the fact that, when dealing with large amounts of data, the traditional tools are usually insufficient and that we must look at the challenges with a different perspective.

Regarding our satisfaction with the delivered code, not much can be said since we did not have access to the answers considered 'more correct'. However, taken in consideration the guidance given by our professor for the average execution time, and looking at a few results shared by other class colleagues, we believe that our solution fulfills all task requirements and surpasses the expectations in terms of performance.

References

1. S. Matos, *IR: Assignment 1*, University of Aveiro, 2019/20.

2. Abstract Base Classes, Python.org, *https://docs.python.org/2/library/abc.html*, (visited in 01/10/2019)

3. Snowball Stem PyStemmer, GitHub.com, *https://github.com/snowballstem/pystemmer*, (visited in 01/10/2019)