

Assignment 2

Filipe Pires [85122], João Alegria [85048]

Information Retrieval

Department of Electronics, Telecommunications and Informatics

University of Aveiro

November 2, 2019

Introduction

This report follows the work delivered for the first assignment of the discipline of 'Information Retrieval' and describes both the updates made on the initial solution and the development of new features related to the indexing of text corpus.

We include the updates done on each class developed and on the respective methods, as well as the redesign of our class diagram. We also provide the instructions on how to run our code.

An explanation on the implementation of the SPIMI approach for the indexing process that considers memory limitations is presented in this report. Changes to the output file (the actual index) format are also explained.

Along with the description of the solution, we also answer to a few questions proposed for the assignment (1). All code and documentation is present in our public GitHub at <https://github.com/joao-alegria/RI>.

1. Updates to Assignment 1

Once concluded the period of time dedicated for the development of the work proposed for the first assignment, a few problems with our delivery were detected and we found that their solution, although simple, was important to be presented in this next assignment.

In point d) of the 4th task of assignment 1, it was asked for us to present the ten terms with highest document frequency according to the index produced by our program using both tokenizer implementations. Our mistake was that we presented the ten terms with highest collection frequency. To make up for this error, we decided to present here the ten terms with highest document frequency, as well as those with highest collection frequency and those with highest term frequency. The difference between these 3 statistical properties are presented below:

- Document Frequency - the number of documents in a collection where a given token occurs.
- Term Frequency - the number of occurrences of a given token in a document.
- Collection Frequency - the total number of occurrences of a given token in the collection.

The 10 terms with highest document frequency (for both tokenizers) are the following:

Simple:

Terms: 'and' (1739514), 'the' (1589467), 'with' (598916),
'for' (585442), 'from' (229502), 'patients' (226388),
'human' (211802), 'cell' (171278), 'cells' (167779),
'study' (167384).

Complex:

Terms: 'cell' (323874), 'patient' (284668), 'effect' (279130),
'human' (226069), 'studi' (215412), 'activ' (192730),
'use' (176649), 'protein' (175199), 'rat' (171824),
'diseas' (165580).

The 10 terms with highest term frequency (for both tokenizers) are the following:

Simple:

Terms: ['the'] (14), ['alpha', 'ucdeq'] (12), ['mcm'] (11),
['and', 'eta', 'mcm', 'nov', 'nvheq', 'sorbitan', 'tky'] (10).

Complex:

Terms: ['alpha'] (11), ['eta', 'nov', 'sarcosin', 'sorbitan'] (10),
['edta', 'failur', 'kinas', 'manp'] (9), ['beta', 'buttiiauxella',
'hla', 'mycoplasma', 'placent', 'pseudomona', 'silic', 'subsp',
'tick', 'val'] (8).

The 10 terms with highest collection frequency (for both tokenizers) are the following:

Simple:

```
Terms: 'and' (2044099), 'the' (2033707), 'with' (633062),  
'for' (608296), 'from' (234101), 'patients' (228824),  
'human' (217457), 'cell' (184034), 'cells' (174998),  
'study' (170588)
```

Complex:

```
Terms: 'cell' (359036), 'patient' (289130), 'effect' (281633),  
'human' (232049), 'studi' (219451), 'activ' (202904),  
'protein' (188952), 'use' (178166), 'rat' (173397),  
'diseas' (169893)
```

Let us now take a look at the updates made for the purpose of this second assignment.

The first update we made to our code was on the content of the index files produced by our indexer. Initially it would write to a text file a token per line followed by a comma (',') and all document IDs of the documents that contained the respective token, paired with the respective frequency (with a ':' between each docID and frequency value). For this assignment, we altered this format in order to store the position of each occurrence and the term frequency weight (tfw) and inverse document frequency (idf) of each term, following the given format:

```
token:idf, docID:tfw:pos1,pos2,...; docID:tfw:...; ...
```

Actually, we implemented this in such way that it is possible for the user to choose from 4 different combinations of output formats, varying the presentation of either the frequency or the weight and the presentation or not of the positions of the terms in each document. Although this was not specifically asked in the assignment instructions, it answers its requirements and gives the user a wider range of usage alternatives.

Regarding the structure of our abstract classes, a few changes had to be considered to better face the challenges of the second assignment. These, however, were not very relevant in terms of difficulty or time consumption.

The capacity of regulating the amount of memory used during the program's execution was a more complex task, with several aspects that required close attention. Our work on the implementation of a SPIMI approach is explained in detail further ahead, but it is good to mention that a few updates to the *main()* function had to be made in order for this and all other changes to be supported by the new version of the program. Still on the updates to the initial phases of the program execution, it is important to mention that we altered the way we process the arguments passed to the scripts when running them. The changes were made to employ the standard form of dealing with script parameters, with the help of the `getopt` Python library (2).

2. Architecture & Memory Usage

The remaining of the report supposes that the reader has knowledge about what was presented on the first delivery. We will only explain the structure of new components not yet seen by the reader and the updates made on the class diagram.

In this chapter we also get into deeper detail on how is memory usage considered, as well as on how are the new attributes such as the term weights calculated and stored. Figure 1 shows the new class diagram, with a considerable increase in complexity compared to the initial version.

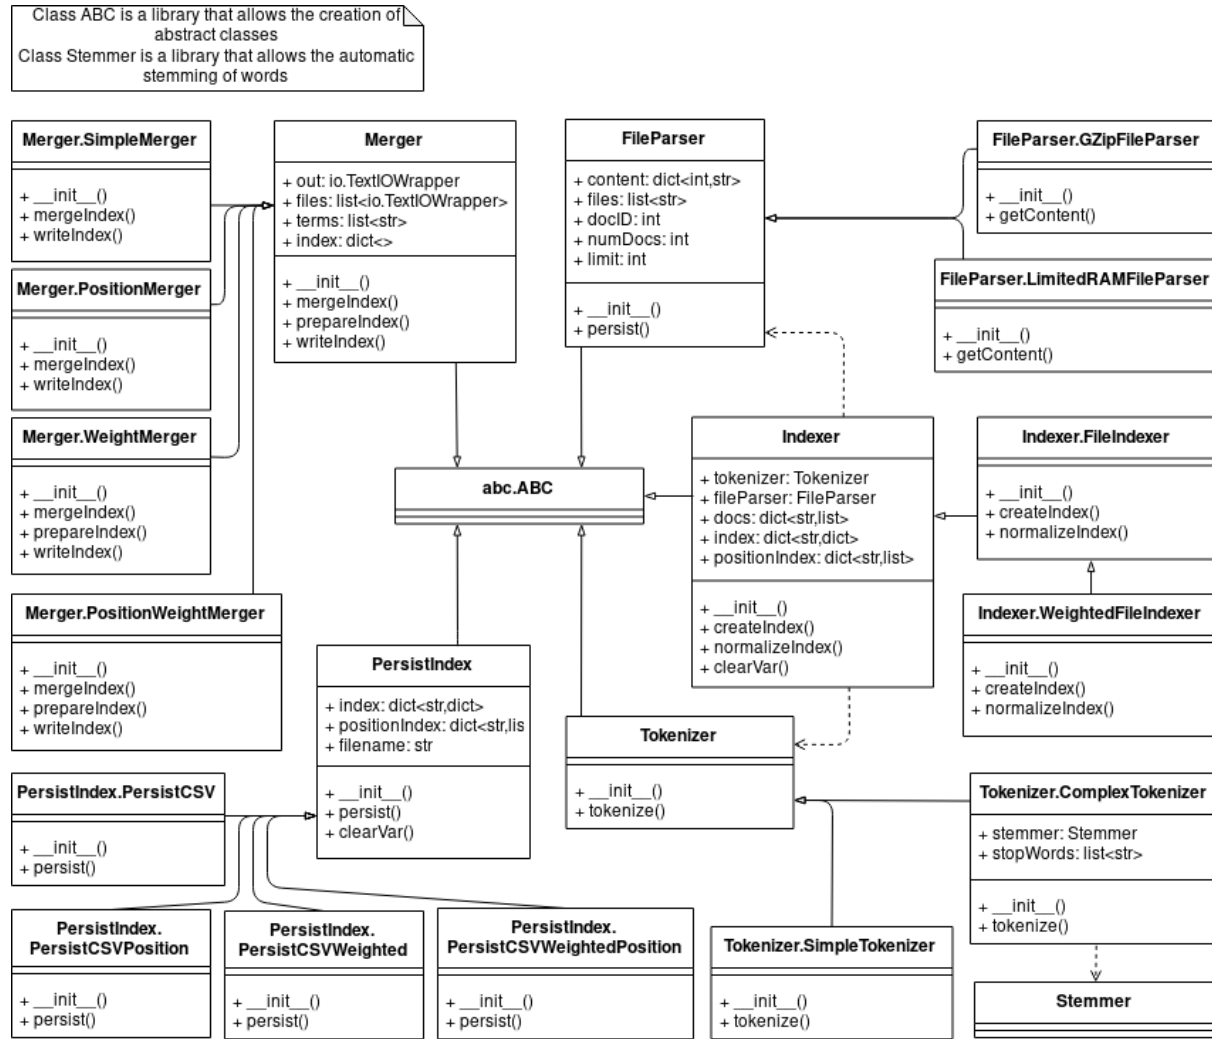


Figure 1: Program's class diagram.

First of all, we have the new class derived from `FileParser` with the name `LimitedRAMFileParser`. This class does the same as its "sister" `GZipFileParser`, except that `getContent()` is now called 1 time per document and returns the information relative to that document only. Although this makes the reading process slower, it also drastically reduces the amount of RAM required during execution.

The `WeightedFileIndexer` is a new class derived from our `FileIndexer`. This class inherits the method *createIndex()* from our previous implementation and makes use of the method *normalizeIndex()* to calculate the term frequencies.

.....

The `Tokenizer` implementations were the only classes that suffered no changes nor received no new subclasses. Our solution served perfectly for the tokenization process whether the program is executed with a memory limitation or not.

For the `PersistIndex` class, we developed 4 new classes derived from it.

.....

// explain the implementation of SPIMI

3. Index Merging

//

4. Discussion

To test the capabilities of the developed software, we indexed the same 2 large compressed files used on the previous deliver, with the names `2004_TREC_ASCII_MEDLINE_1.gz` and `2004_TREC_ASCII_MEDLINE_2.gz`. Their structure, as collections (corpus) of documents, remained the same. In this chapter we discuss our implementation's efficiency over these files according to the following measures:

- a) What is the total indexing time and final index size on disk?
- b) What is the maximum amount of memory used during indexing?

//

Conclusions

After completing the assignment, we drew a few conclusions regarding our solutions and the whole concept of memory management and its balancing with good efficiency.

//

References

1. S. Matos, *IR: Assignment 2*, University of Aveiro, 2019/20.
2. GetOpt Library - C-Style Parser for Command Line Options, Python.org
<https://docs.python.org/2/library/getopt.html>, (visited in 25/10/2019)