Assignment 2

Filipe Pires [85122], Joao Alegria [85048]

Information Retrieval
Department of Electronics, Telecommunications and Informatics
University of Aveiro

November 11, 2019

Introduction

This report was written as a follow-up work of the delivery of the first assignment for the discipline of 'Information Retrieval' and describes both the updates made on the initial solution and the development of new features related to the indexing of text corpus.

We include the updates done on each class developed and on the respective methods, as well as the redesign of our class diagram. We also provide the instructions on how to run our code.

An explanation on the implementation of the SPIMI approach for the indexing process that considers memory limitations is presented in this report. Changes to the output file (the actual index) format are also explained.

Along with the description of the solution, we also answer to a few questions proposed for the assignment (*1*). All code and documentation is present in our public GitHub project at `https://github.com/joao-alegria/RI`.

1. Updates to Assignment 1

Once concluded the period of time dedicated for the development of the work proposed for the first assignment, some mistakes on our delivery were detected and we found that their solution, although simple, was important to be presented in this next assignment.

In point d) of the 4th task of assignment 1, it was asked for us to present the ten terms with highest document frequency according to the index produced by our program using both tokenizer implementations. Our mistake was that we presented the ten terms with highest collection frequency. To make up for this error, we decided to present here the ten terms with highest document frequency, as well as those with highest collection frequency and those with highest term frequency. The difference between these 3 statistical properties are presented below:

- Document Frequency - the number of documents in a collection where a given token occurs.
- Term Frequency - the number of occurrences of a given token in a document.
- Collection Frequency - the total number of occurrences of a given token in the collection.

The 10 terms with highest document frequency (for both tokenizers) are the following:

```
Simple:
Terms: 'and' (1739514), 'the' (1589467), 'with' (598916),
'for' (585442), 'from' (229502), 'patients' (226388),
'human' (211802), 'cell' (171278), 'cells' (167779),
'study' (167384).

Complex:
Terms: 'cell' (323874), 'patient' (284668), 'effect' (279130),
'human' (226069), 'studi' (215412), 'activ' (192730),
'use' (176649), 'protein' (175199), 'rat' (171824),
'diseas' (165580).
```

The 10 terms with highest term frequency (for both tokenizers) are the following:

```
Simple:
Terms: ['the'] (14), ['alpha','ucdeq'] (12), ['mcm'] (11),
['and','eta','mcma','nov','nvheq','sorbitan','tky'] (10).

Complex:
Terms: ['alpha'] (11), ['eta','nov','sarcosin','sorbitan'] (10),
['edta','failur','kinas','manp'] (9), ['beta','buttiauxella',
'hla','mycoplasma','placent','pseudomona', 'silic','subsp',
'tick','val'] (8).
```

The 10 terms with highest collection frequency (for both tokenizers) are the following:

```
Simple:
Terms: 'and'(2044099), 'the'(2033707), 'with'(633062),
'for'(608296), 'from'(234101), 'patients'(228824),
'human'(217457), 'cell'(184034), 'cells'(174998),
'study'(170588)

Complex:
Terms: 'cell'(359036), 'patient'(289130), 'effect'(281633),
'human'(232049), 'studi'(219451), 'activ'(202904),
'protein'(188952), 'use'(178166), 'rat'(173397),
'diseas'(169893)
```

Let us now take a look at the updates made for the purpose of this second assignment.

The first update we made to our code was on the content of the index files produced by our indexer. Initially it would write to a text file a token per line followed by a comma (',') and all document IDs of the documents that contained the respective token, paired with the respective frequency (with a ':' between each docID and frequency value). For this assignment, we altered the format in order be possible to store the position of each occurrence and the term frequency weight (tfw) and inverse document frequency (idf) of each term, following the given format:

```
token:idf, docID:tfw:pos1,pos2,...; docID:tfw:...; ...
```

Actually, we implemented this in such way that it is possible for the user to choose from 4 different combinations of output formats, varying the presentation of either the frequency or the weight and the presentation or not of the positions of the terms in each document. Although this was not specifically asked in the assignment instructions, it answers its requirements and gives the user a wider range of usage alternatives.

Regarding the structure of our abstract classes, a few changes had to be considered to better face the challenges of the second assignment. These, however, were not very relevant in terms of difficulty or time consumption.

The capacity of regulating the amount of memory used during the program's execution was a more complex task, with several aspects that required close attention. Our work on the implementation of a SPIMI approach is explained in detail further ahead, but it is good to mention that a few updates to the *main()* function had to be made in order for this and all other changes to be supported by the new version of the program. Still on the updates to the initial phases of the program execution, it is important to mention that we altered the way we process the arguments passed to the scripts when running them. The changes were made to employ the standard form of dealing with script parameters, with the help of the `getopt` Python library (*2*).

2. Architecture & Memory Usage

The remaining of the report supposes that the reader has knowledge about what was presented on the first delivery. We will only explain the structure of new components not yet seen by the reader and the updates made on the class diagram.

In this chapter we also get into deeper detail on how is memory usage considered, as well as on how are the new attributes such as the term weights calculated and stored. Figure 1 shows the new class diagram, with a considerable increase in complexity compared to the initial version.

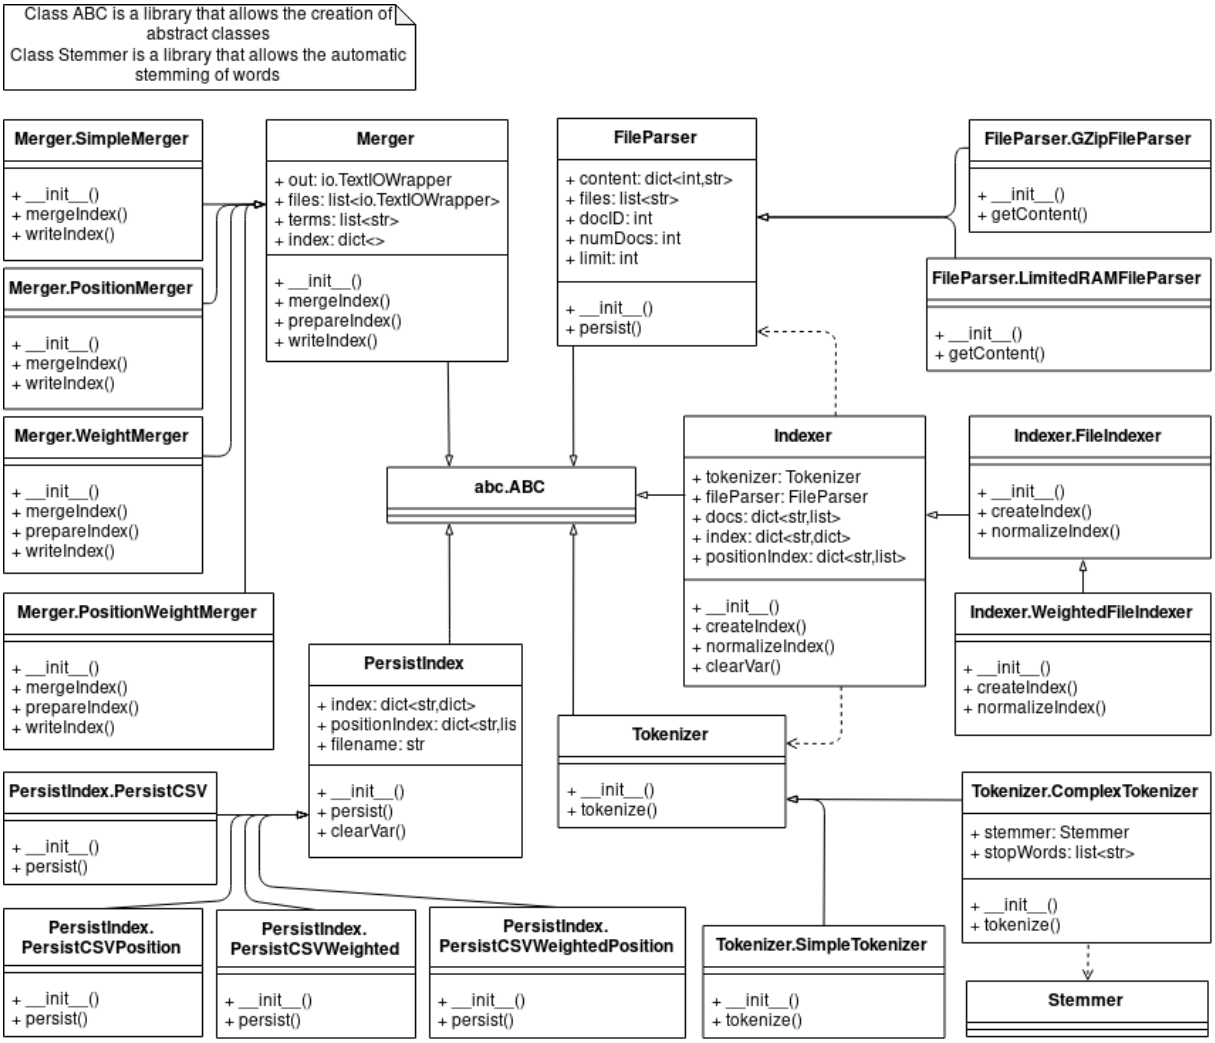

Figure 1: Program's class diagram.

First of all, we have the new class derived from `FileParser` with the name `LimitedRAMFileParser`. This class does the same as its "sister" `GZipFileParser`, except that *getContent()* is now called 1 time per document and returns the information relative to that document only. Although this makes the reading process slower, it allowes the user to have a improved control over the RAM consumption.

The `WeightedFileIndexer` is a new class derived from our `FileIndexer`. This class inherits the method *createIndex()* from our previous implementation and makes use of the method *normalizeIndex()* to calculate the term frequencies. We call it a normalization as it is a transformation of the produced index to standardize its values according to a set of rules. We calculate the term frequencies following the logarithmic formula variant and the process of normalization is based on the cosine formula variant, both presented bellow.

$$tf_{t,d} = 1 + log(tf_{t,d}) \qquad W = \frac{1}{\sqrt{w_1^2 + w_2^2 + ... + w_M^2}}$$

The `Tokenizer` implementations were the only classes that suffered no changes nor recieved no new subclasses. Our solution served perfectly for the tokenization process whether the program is executed with a memory limitation or not.

For the `PersistIndex` class, we developed 3 new classes derived from it - one for each new output format. Following are small portions of indexes generated by our solution in the formats supported by the program, for analysis purposes(all with the complex tokenizer).

Simple Output(in output/complexSample.txt):

...
acetyl,133:1
acid,54:1,55:1,65:1,72:1,81:1,86:1,91:1,104:1,112:1,116:1,126:1,159:1,171:1
...

Positions Output(in output/complexP):

...
acetyl;133:1:4
acid;54:1:10;55:1:7;65:1:4;72:1:4;81:1:2;86:1:8;91:1:4;104:1:10;112:1:4;116:1:7;126:1:8;159:1:4;171:1:7
...

Weights Output(in output/complexW):

...
acetyl:2.35;133:1.0
acid:1.24;54:0.28;55:0.28;65:0.28;72:0.28;81:0.28;86:0.28;91:0.28;104:0.28;112:0.28;
116:0.28;126:0.28;159:0.28;171:0.28
...

Weights+Positions Output(in output/complexWP):

...

acetyl:2.35;133:1.0:4

acid:1.24;54:0.28:10;55:0.28:7;65:0.28:4;72:0.28:4;81:0.28:2;86:0.28:8;91:0.28:4;
104:0.28:10;112:0.28:4;116:0.28:7;126:0.28:8;159:0.28:4;171:0.28:7

...

We now proceed to explaining how was the memory considered in this second assignment and, more specifically, what is the SPIMI strategy already mentioned in this document. SPIMI stands for Single-Pass In-Memory Indexing, and its key ideas are that several pseudo-indexes (index approximations) are generated, one per block of information (where, in our case, a block represents the amount of information capable of being stored in memory alone), accumulating postings (details about each term) in postings lists as they occur. Then, these block-indexes are merged with a n-way merge, for efficiency purposes, into a unique index. This final index, however, is stored in one or more output files, according to the memory restrictions applied to the program.

Our solution has the following execution flow: while there is still memory available, the program parses one document, tokenizes it and adds it to an internal index; if the memory limit is reached, this internal index is then normalized and persisted to disk according to the chosen format, and finally the memory is cleared and the cicle continues. Once the entire corpus has been processed, if needed, the merger is created and the n-way merge is executed to produce the final index. The memory used by the process responsible for the program's execution is assured never to exceed the memory limitation applied by the user.

3. Index Merging

We have not yet talked about the `Merger` class, although it has already been mentioned that a merging process takes place in this second assignment. This chapter is dedicated to this new module of our solution and the decisions made during its implementation.

How exactly does the n-way merge work? The most straightforward solution would be to merge index files in pairs, by directly comparing each line and merging each into one. This process would repeat until all files were reduced to one. The problem with this alternative is that, after each pair merging process, the newly created index file will have to be processed again - this adds a redundancy that severly slows the execution. This is where the n-way merging fits in. To solve the redundancy problem, we iterate over all index files (n means the number of files) simultaneously. This way, we guarantee that each line is only processed once.

The way we implemented the n-way merge works as follows. With the aid of a temporary array, we store the latest line of each index file, one at a time; when this occurs for all files, we calculate the minimum term by finding the minimum (alphabetically) of the auxiliary array. If the minimum term appears more than once in the array, the merging of the respective postings lists happens, followed by storing those values in an internal (in-memory) index. The positions in the temporary array where the merging of the term occurred are replaced by an empty value. New lines are only inserted where the array contains these empty values and they are fetched from the respective index file. Now the process repeats itself, constantly merging the terms that occur in more than one index file, and storing the merged term in the internal index. Once this in-memory index reaches a point where there is no more free memory available, its content is then written to disk and the memory is freed so that the process can continue with the remaining terms.

4. Discussion

To test the capabilities of the developed software, we indexed the same 2 large compressed files used on the previous deliver, with the names `2004_TREC_ASCII_MEDLINE_1.gz` and `2004_TREC_ASCII_MEDLINE_2.gz`. Their structure, as collections (corpus) of documents, remained the same. In this chapter we discuss our implementation's efficiency over these files according to the following measures:

a) What is the total indexing time and final index size on disk?

b) What is the maximum amount of memory used during indexing?

Although running the program works very similarly to the previous version, we present here the updated format. The file `CreateIndex.py` serves as the pipeline creator and is ther command to be executed in order to indexing all documents passed as arguments, and it has the following execution options:

```
$ python3 CreateIndex.py [-h] [-p] [-w] [-o outputFolder]
  [-l limit] [-t tokenizer] [-r limitRAM] inputFolder
```

Here, `-h` is the option that presents the manual for the command usage. Options `-p` and `-w`, if present, tell the program to process the positions and weights of the terms in the index, respectively. Options `-o` and `-l` allow the definition of the output folder's name and of the limit for the number of lines to be processed in each input file. Option `-t` makes possible for the user to choose the type of tokenizer to be used, and the alternatives are: 'simple' for the use of the `SimpleTokenizer` class, and 'complex' for the `ComplexTokenizer` class. Option `-r` allows the user to define the maximum amount of memory that can be used by the process running the program. The previous arguments are all optional and the actual values for these arguments must appear right after the respective options. The final argument of the command must be the name of the folder that contains the files to be indexed. This argument was changed since last assignment as we considered more user-friendly to identify the target folder rather than each one of the input files (this is more noticeable as the number of files increases). Other argument that changed was the output option, transitioning from file name to folder name since now it's possible to create more than one file as the output index and to maintain a user-friendly application we decided to make this adjustment.

The final version of our solution was executed according to the following format, as requested on the question 2 (the input folder contained the compressed files previously mentioned):

```
$ command time -v python3 CreateIndex.py -o ../index -w
  -t complex -r 0.5 ../input
```

The results are here presented. The total indexing time, given by the command ”command time -v” was of 60 minutes . The final index size on disk is calculated by summing the size on disk of the various sub-indexes generated by the merger. The result was of 491.4 Mb .

The maximum amount of memory used during indexing was also given by the ”command time -v” command and had a value of 462 Mb , meaning that the memory limit given of 500Mb was indeed not surpassed.

In relation to the question 3, the final results were obtained by running the following command:

```
$ command time -v python3 CreateIndex.py -o ../index -wp
  -t complex -r 0.5 ../input
```

The total indexing time was of 61 minutes. The final index size on disk was 577.3 Mb . Finnaly the maximum amount of memory used during indexing was 487 Mb.

Taking a closer look at the results, the first fact that takes our attention is the execution time. When we compare it to the initial solution (of the first assignment), we see an increase in time of 500 %. This is explained by the fact that, not only the reading process becomes slower as the number of accesses done to the input files are increased, but also the number of writes to disk also increase as well as the number of files we write to. All these I/O interactions, i.e. accesses to disk, are very time consuming and are a good example of the well-known bottleneck problem of high resource-demanding tasks. The advantage of these interactions lies on the reduction of other resource usages - more specifically, of RAM usage. By indexing large files with this method, we guarantee that the program can process corpus that far exceed the hardware resources available. The user now gains power to balance resource usage according to his / her needs and available hardware.

The second observation we find important to refer regards the final index storage strategy. We decided to store the outcome in several partial index files since we took into account that, in the future, the aim will be to retrieve information from them through queries and so these partial index files will allow for a faster retrieval of results.

Conclusions

After completing the assignment, we drew a few conclusions regarding our solutions and the whole concept of memory management and its balancing with good efficiency.

The biggest challenge we faced was definetly the merging process. As we wished to come up with a solution of our own, rather than using external code or libraries, the planning and execution took a considerable amount of time and effort. In addition, testing was also not very much straightforward, as the results needed to be manually checked. The delivered code, although time consuming during its execution, ensures that no term is lost, no attribute wrongly calculated and the results are exactly the same compared to the outcome of the execution without memory limitation (i.e. without the merging of intermediate indexes).

From this assignment, we take the importance of being able to control the tradeoff between indexing time and memory usage. Although execution times greatly increases once memory restrictions are considered, we saw a great value in being able to constrain the resources consumed by the process in order to allow for the operating system to still offer sufficient memory for other tasks.

When it comes to teamwork, we follow a balanced work distribution, usually dedicating the time for the assignment in compatible schedules. This helps keeping up-to-date with what each is doing and speeds up problem solving.

References

1. S. Matos, *IR: Assignment 2*, University of Aveiro, 2019/20.

2. GetOpt Library - C-Style Parser for Command Line Options, Python.org *https://docs.python.org/2/library/getopt.html*, (visited in 25/10/2019)